



Analyzing ELF Binaries with Malformed Headers Part 1 - Emulating Tiny Programs



JULIAN - SEPTEMBER 17, 2019

A simple but often effective method for complicating or preventing analysis of an ELF binary by many common tools (`gdb`, `readelf`, `pyelftools`, etc) is mangling, damaging or otherwise manipulating values in the ELF header such that the tool parsing the header does so incorrectly, perhaps even causing the tool to fail or crash. Common techniques include overlapping the ELF header with the program header table and writing non-standard values to ELF header fields that are not needed for composing the process image of the binary in memory. In addition to some programs designed for criminal purposes (e.g. the “mumblehard” family of malware programs), a few code-golf- and proof-of-concept-type programs have been created that employ these techniques. Examples of such programs include Brian Raiter’s “teensy” files and @netspooky’s “golfclub” programs. In this post, it will be demonstrated how emulation can be used to trace the execution of these types of binaries.

Overview

The following will be discussed:

- how header mangling works as an anti-analysis technique
- how to use the Unicorn Engine to analyze minimalist binaries with malformed headers

Tools:

- Capstone Engine
- Unicorn Engine
- Python3

Malformed ELF Headers

This technique has already been covered in depth elsewhere^{[1][2][3][4][5][6]}, so the discussion here will be brief. The main reason mangling the ELF header works to complicate analysis is that even though only a specific subset of the fields in the ELF header are read by the kernel when loading the program into memory, most ELF parsers do not parse the ELF header the way the kernel loader does and thus are prone to malfunction when reading unexpected or garbage values in these extraneous (from the perspective of loading) fields. The most typical examples of this are `gdb`, `objdump` and the rest of the `libbfd`-based binutils tools, which will not even read an object file unless its section information is present and intact.

The specially-crafted minimalist ELF programs - those that push the limits of the least number of bytes a file can consist of and still execute successfully - take advantage of the fact that not all ELF header fields are needed for loading and executing the program and can therefore be used to contain code or other non-standard values; as a case in point, the entry points of these programs often lie *inside* their ELF header. On the one hand, even though complicating analysis is not an explicit goal of their design, these programs serve to highlight the limitations of many common tools designed to work with the ELF format. On the other hand, since these minimalist binaries typically contain such little code, using fully-featured debuggers and other tools of this class for analysis would actually be overkill; one may have a good laugh about how NSA’s Ghidra cannot properly load their tiny ELF file, but attempting use such a tool to analyze an extremely minimalist binary is akin to trying to shoot down a fruitfly with a railgun - heavyweight frameworks packaged with disassemblers plus debuggers and/or decompilers are unsuitable and unnecessary for analyzing the runtime behavior of

executables literally 45 or 62 bytes in size. If there are 10 bytes of code in a program, does it make sense to try to load it into a decompiler? Probably not. A simple script emulating the execution of these programs may be a more appropriate approach.

muppetlabs' tiny-i386: 45 bytes total, 7 bytes of code

This is the “Tiny” program from [A Whirlwind Tutorial on Creating Really Teensy ELF Executables for Linux](#). This program, as well as the rest of the “Teensy” ELF files can be downloaded from the [muppetlabs site](#).

The approach taken here to analyzing this file is as follows:

- attempting analysis with `readelf`, `gdb` and `r2`
- looking at the source code
- emulation

Using Standard Tools to Parse the file

It should be noted at the outset that the binary can be loaded and executed without any problems:

```
$ strace ./tiny-i386
execve("./tiny-i386", ["/tiny-i386"], 0x7ffc2a421f60 /* 52 vars */) = 0
strace: [ Process PID=30049 runs in 32 bit mode. ]
exit(42)                                = ?
+++ exited with 42 +++
```

However when `readelf` is used to try to read the program's ELF header, it fails:

```
$ readelf -h tiny-i386
readelf: Error: tiny-i386: Failed to read file header
```

`gdb` fails to recognize that it is indeed an ELF file:

```
$ gdb -q tiny-i386
GEF for linux ready, type `gef' to start, `gef config' to configure
80 commands loaded for GDB 8.1.0.20180409-git using Python engine 3.6
"home/reversing/tiny-i386": not in executable format: File format not recognized
gef> info file
gef> run
Starting program:
No executable file specified.
Use the "file" or "exec-file" command.
gef>
```

In a pleasant surprise, we are able to debug and disassemble the code using `r2`:

```
$ r2 -d tiny-i386
Process with PID 29855 started...
```

```

= attach 29855 29855
bin.baddr 0x00010000
Using 0x10000
Warning: Cannot initialize program headers
Warning: Cannot initialize section headers
Warning: Cannot initialize strings table
Warning: Cannot initialize dynamic strings
Warning: Cannot initialize dynamic section
Warning: read (init_offset)
asm.bits 32
[0x00010020]> ds
[0x00010020]> ds
[0x00010020]> ds
[0x00010020]> ds
child exited with status 42

==> Process finished

Stepping failed!
Step failed
[0x00010020]> pd 10
      0x00010020      b32a      mov bl, 0x2a      ; ebx
      0x00010022      31c0      xor eax, eax
      0x00010024      40        inc eax
      ;-- eip:
      0x00010025      cd80      int 0x80      <----- exec
      0x00010027      003400    add byte [eax + eax], dh
      0x0001002a      2000      and byte [eax], al
      0x0001002c      ~ 0100    add dword [eax], eax
      ;-- section_end.ehdr:
      0x0001002d      0000      add byte [eax], al
      0x0001002f      0000      add byte [eax], al
      0x00010031      0000      add byte [eax], al
[0x00010020]>

```

However, when radare2 is used to parse the binary, some of the field values look strange:

```

$ r2 -nn tiny-i386
[0x00000000]> pf.elf_header @ elf_header
      ident : 0x00000000 = .ELF.
      type  : 0x00000010 = type (enum elf_type) = 0x2 ; ET_EXEC
      machine : 0x00000012 = machine (enum elf_machine) = 0x3 ; EM_386
      version : 0x00000014 = 0x00010020
      entry  : 0x00000018 = 0x00010020
      phoff  : 0x0000001c = 0x00000004
      shoff  : 0x00000020 = 0xc0312ab3

```

```

    flags : 0x00000024 = 0x0080cd40
    ehsize : 0x00000028 = 0x0034
    phentsize : 0x0000002a = 0x0020
    phnum : 0x0000002c = 0xff01
    shentsize : 0x0000002e = 0xffff
    shnum : 0x00000030 = 0xffff
    shstrndx : 0x00000032 = 0xffff
[0x00000000]>

```

There do seem to be quite a few odd-looking values mixed together with ones that appear similar to what we are accustomed to seeing. What is happening here? Examining the source code will help explain some of this output.

A Look at the Source Code

```
; tiny.asm
```

```
BITS 32
```

```

                                org      0x00010000

                                db        0x7F, "ELF"                ; e_ident
                                dd        1                            ; p
                                dd        0                            ; p
                                dd        $$                          ; p
                                dw        2                            ; e_type
                                dw        3                            ; e_machine
                                dd        _start                       ; e_version
                                dd        _start                       ; e_entry
                                dd        4                            ; e_phoff
_start:
                                mov       bl, 42                      ; e_shoff
                                xor       eax, eax
                                inc       eax                          ; e_flags
                                int       0x80
                                db        0
                                dw        0x34                      ; e_ehsize
                                dw        0x20                      ; e_phentsize
                                db        1                          ; e_phnum
                                ; e_shentsize
                                ; e_shnum
                                ; e_shstrndx

```

```
filesize      equ      $ - $$
```

A few observations:

- the program header overlaps with the ELF header
- the entry point is inside the ELF header
 - The implication is that there is executable code inside the header
- the fields having to do with sections are empty
 - it is actually more precise to say that since the file is 45 bytes in size but the ELF header of a 32-bit binary should be 52 bytes in total, those fields are simply not there.

As it turns out, the subset of fields that must contain correct values in order to be loaded by the kernel consists of the following:

- The first 4 bytes of `e_ident` which includes:
 - `EI_MAG0 - EI_MAG4`: `0x7f`, `E`, `L`, `F`
- `e_type`
- `e_machine`
- `e_entry`
- `e_phoff`
- `e_phnum`

Summary from “A Whirlwind Tutorial on Creating Really Teensy ELF Executables for Linux” (bolding added):

So: Here’s what is and isn’t essential in the ELF header. The **first four bytes** have to contain the magic number, or else Linux won’t touch it. The other three bytes in the `e_ident` field are not checked, however, which means we have no less than twelve contiguous bytes we can set to anything at all. **`e_type`** has to be set to 2, to indicate an executable, and **`e_machine`** has to be 3, as just noted. `e_version` is, like the version number inside `e_ident`, completely ignored. (Which is sort of understandable, seeing as currently there’s only one version of the ELF standard.) **`e_entry`** naturally has to be valid, since it points to the start of the program. And clearly, **`e_phoff`** needs to contain the correct offset of the program header table in the file, and **`e_phnum`** needs to contain the right number of entries in said table. `e_flags`, however, is documented as being currently unused for Intel, so it should be free for us to reuse. `e_ehsize` is supposed to be used to verify that the ELF header has the expected size, but Linux pays it no mind. `e_phentsize` is likewise for validating the size of the program header table entries. This one was unchecked in older kernels, but now it needs to be set correctly. Everything else in the ELF header is about the section header table, which doesn’t come into play with executable files.

Emulation

Given that the program contains only 7 bytes of instructions and has a malformed header, emulation is a good alternative to heavyweight tools like radare2, IDA, Ghidra, etc. for analyzing/tracing/logging the runtime behavior of this kind of program. The program’s code can be emulated via a small python script that utilizes the [Unicorn Engine](#) (at time of writing, the [Qiling emulation framework](#) is still in alpha and the code is not available). For our purposes right now, it does not matter that the ELF header is malformed, as the only information needed to emulate the binary is its architecture and the file offsets at which to begin and end emulation; this information can be retrieved from a hex dump of the binary without needing to parse the ELF header.

The approach to emulating the `tiny-i386` binary is as follows:

First, retrieve the start and end points for emulation from a hex dump. Then, when writing the script to emulate the program:

- read the file and map it to memory
- set up the stack
- initialize the emulation engine
- implement a hook that allows each executed instruction to be traced and logged to STDOUT
 - a Capstone disassembly engine object will be passed to this hook so that each instruction can be disassembled and its disassembly logged as well
- implement a hook that handles system calls

We know from the source code that the first instruction is `mov bl, 42` and the final instruction is `int 0x80`. We can find these easily in a dump of `tiny-i386`:

```
$ hexdump -C tiny-i386
00000000  7f 45 4c 46 01 00 00 00  00 00 00 00 00 00 01 00  |.ELF.....|
```

```

00000010  02 00 03 00 20 00 01 00  20 00 01 00 04 00 00 00  |.... ... ..|
00000020  b3 2a 31 c0 40 cd 80 00  34 00 20 00 01          |.*1.@...4. ..|
0000002d

```

Narrowing down the output:

```

$ hexdump -C -s 0x20 -n 7 tiny-i386
00000020  b3 2a 31 c0 40 cd 80          |.*1.@...|
00000027

```

There we have it: the offset at which to begin emulation is `0x20` and at which to end is `0x27`.

Since the architecture is already known to us, this is all the information required to emulate the program:

```

1  #!/usr/bin/python3
2
3  from unicorn import *
4  from unicorn.x86_const import *
5  from capstone import *
6  import struct
7
8
9  BASE = 0x100000
10 STACK_ADDR = 0x0
11 STACK_SIZE = 1024 * 1024
12
13 def read(name):
14     with open(name, 'rb') as f:
15         return f.read()
16
17 #https://github.com/unicorn-engine/unicorn/blob/master/bindings/python/shellcode.py
18 # callback for tracing instructions
19 def hook_code(uc, address, size, user_data):
20     instruction = uc.mem_read(address, size) # read this instruction code from mem
21     md = user_data
22     for i in md.disasm(instruction, address):
23         print(">>> Tracing instruction at 0x%x, instruction size = 0x%x, disassembly:"
24
25
26 # callback for tracing Linux interrupt
27 def hook_intr(uc, intno, user_data):
28     # only handle syscall
29     if intno != 0x80:
30         print("got interrupt %x ??? " % intno);
31         uc.emu_stop()

```

```
32         return
33
34     eax = uc.reg_read(UC_X86_REG_EAX)
35     eip = uc.reg_read(UC_X86_REG_EIP)
36
37     print(">>> 0x%x: INTERRUPT: 0x%x, EAX = 0x%x" %(eip, intno, eax))
38
39     uc.emu_stop()
40
41
42
43 def main():
44
45     mu = Uc(UC_ARCH_X86, UC_MODE_32)    # initialize emulation engine class
46     mu.mem_map(BASE, STACK_SIZE)        # allocate space at base address
47     mu.mem_map(STACK_ADDR, STACK_SIZE)  # allocate space for stack
48
49     mu.mem_write(BASE, read("./target_binaries/tiny-i386"))    # write file to memory
50     mu.reg_write(UC_X86_REG_ESP, STACK_ADDR + STACK_SIZE - 1) # initialize stack
51
52     md = Cs(CS_ARCH_X86, CS_MODE_32)    # initialize disassembler engine class
53
54     # add hooks
55     mu.hook_add(UC_HOOK_CODE, hook_code, md)    # pass disassembler engine to hook
56     mu.hook_add(UC_HOOK_INTR, hook_intr)
57
58     mu.emu_start(BASE + 0x20, BASE + 0x27)
59
60     print(">>> Emulation Complete.")
61
62 if __name__ == "__main__":
63     main()
```

emulate_tiny-i386.py hosted with ❤ by GitHub [view raw](#)

When executed, we get a trace + disassembly:

```
$ ./emulate_tiny-i386.py
>>> Tracing instruction at 0x100020, instruction size = 0x2, disassembly:    mov
>>> Tracing instruction at 0x100022, instruction size = 0x2, disassembly:    xor
>>> Tracing instruction at 0x100024, instruction size = 0x1, disassembly:    inc
>>> Tracing instruction at 0x100025, instruction size = 0x2, disassembly:    int
>>> 0x100025: INTERRUPT: 0x80, EAX = 0x1
>>> Emulation Complete.
```

Looks good. No debugger needed.

netspooky's bye: 84 bytes total, 23 bytes of code

An advantage of emulation over debugging is that the emulated instructions (should) have no effect on the host system. Even if the program being emulated contains code that could potentially damage the system it runs on, its instructions are not actually being executed by the CPU, so emulation poses much less risk than debugging (unless there is some way to escape from the emulator, e.g. [QEMU VM escape](#)). This is useful for analyzing viruses, crimeware, etc. and in this particular case [@Netspooky's](#) `bye` binary, which executes the `reboot` `syscall` with the `LINUX_REBOOT_CMD_POWER_OFF` argument:

On a desktop system, this binary will shut down your computer abruptly. There are some potential side effects from a shutdown like this, but personally I haven't experienced any issues with it. However, on a VPS, this specific `syscall` proves to be a bit of a problem. Since the virtual machine doesn't actually have any of its own physical hardware (it's either virtualized or shared with the host), the power button on a VPS isn't really a thing. By executing a `syscall` that effectively "shuts off the power" to the operating system, this puts the VM in an unknown state. So far, whenever this is run on a VPS, it seemingly wipes out the entire instance.

Obviously it is advantageous to be able to analyze the runtime behavior of such a program without having to actually load it into memory and execute since we do not want our machine to be shut down, and in a way that may potentially damage the system at that. The script used to analyze `tiny-i386` can be modified to support emulation of x86-64 code and of the `reboot` `syscall`. The same approach will be followed as before, with minor adjustments.

Before we begin, however, we can first try to read the file's ELF header with `readelf`, take a look at the source code, and then disassemble its code with Capstone to get a sense of what to expect from emulation.

Parsing the Header with readelf

```
$ readelf -h bye
ELF Header:
  Magic:   7f 45 4c 46 ba dc fe 21 43 be 69 19 12 28 eb 3c
  Class:                               <unknown: ba>
  Data:                                   <unknown: dc>
  Version:                             254 <unknown: %lx>
  OS/ABI:                               <unknown: 21>
  ABI Version:                          67
  Type:                                 EXEC (Executable file)
  Machine:                              Advanced Micro Devices X86-64
  Version:                              0x1
  Entry point address:                  0x4
  Start of program headers:             1 (bytes into file)
  Start of section headers:             28 (bytes into file)
  Flags:                                0x0
  Size of this header:                  0 (bytes)
  Size of program headers:              0 (bytes)
  Number of program headers:            0
  Size of section headers:              0 (bytes)
  Number of section headers:            1
  Section header string table index: 0
readelf: Warning: possibly corrupt ELF header - it has a non-zero program header offset
```


The ELF header is clearly malformed. At least we can see the entry point is at offset 0x4.

The Source Code

1	; 84 byte LINUX_REBOOT_CMD_POWER_OFF Binary Golf			
2	BITS 64			
3	org 0x100000000			
4	;-----+-----+-----+-----			
5	; CODE LISTING	OFFS	ASSEMBLY	CODE COMMENT
6	;-----+-----+-----+-----			
7	db 0x7F, "ELF"	; 0x0	7f454c46	PROTIP: Can use magic as a constant ;)
8	_start:	;----- ----- -----		
9	mov edx, 0x4321fedc	; 0x04	badcfe2143	Moving magic values...
10	mov esi, 0x28121969	; 0x09	be69191228	into their respective places
11	jmp short reeb	; 0x0E	eb3c	Short jump down to @x4c
12	dw 2	; 0x10	0200	
13	dw 0x3e	; 0x12	3e00	
14	dd 1	; 0x14	01000000	
15	dd _start - \$\$; 0x18	04000000	
16	phdr:	;----- ----- -----		
17	dd 1	; 0x1C	01000000	
18	dd phdr - \$\$; 0x20	1c000000	
19	dd 0	; 0x24	00000000	
20	dd 0	; 0x28	00000000	
21	dq \$\$; 0x2C	00000000	
22		; 0x30	01000000	
23	dw 0x40	; 0x34	4000	
24	dw 0x38	; 0x36	3800	
25	dw 1	; 0x38	0100	
26	dw 2	; 0x3A	0200	
27	cya:	;----- ----- -----		
28	mov al, 0xa9	; 0x3C	b0a9	Load syscall
29	syscall	; 0x3E	0f05	Execute syscall
30	dd 0	; 0x40	00000000	Filler, should try to keep as all 0's
31	mov al, 0xa9	; 0x44	b0a9	Load syscall
32	syscall	; 0x46	0f05	Execute syscall
33	dd 0	; 0x48	00000000	Filler, should try to keep as all 0's
34	reeb:	;----- ----- -----		
35	mov edi, 0xf0000000	; 0x4C	b0000000	Load magic "LINUX_REBOOT_CMD_POWER_OFF"
36	jmp short cya	; 0x51	ebe9	Short jmp back to e_shnum/p_filesz @0x3C
37	nop	; 0x53	90	Filler, could use this byte for code.
38	;-----+-----+-----+-----			
39	; Note that we are overlaying the ELF Header with the program headers.			
40	; You have 12 bytes minus your short jump from 0x4-0x10 to store code			
41	; Then you have 8 bytes within the program headers at 0x4c for more			

```

42 ; code, plus e_shentsize and the lower bytes of p_filesz + p_memsz for
43 ; storage and code if you stay within the bounds - still testing.
44 ;
45 ;     LINUX_REBOOT_CMD_POWER_OFF
46 ;         (RB_POWER_OFF, 0x4321fedc; since Linux 2.1.30). The message
47 ;         "Power down." is printed, the system is stopped, and all power
48 ;         is removed from the system, if possible. If not preceded by a
49 ;         sync(2), data will be lost.
50 ; [ Compile ]
51 ; nasm -f bin -o bye bye.nasm
52 ;
53 ; One Liner
54 ; base64 -d <<< f0VMRrrc/iFDvmkZEijrPAIAPgABAAAABAAAAEAAAAcAAAAAAAAAAAAAAAAAAAAQAAA
55 ;
56 ; Syscall reference: http://man7.org/linux/man-pages/man2/reboot.2.html

```

bye.asm hosted with ♥ by GitHub

[view raw](#)

```

1 ; [ Full breakdown ]
2 ; --- Elf Header
3 ; Offset # Value Purpose
4 ; 0-3 A 7f454c46 Magic number - 0x7F, then 'ELF' in ASCII
5 ; 4 B ba 1 = 32 bit, 2 = 64 bit
6 ; 5 C dc 1 = little endian, 2 = big endian
7 ; 6 D fe ELF Version
8 ; 7 E 21 OS ABI - usually 0 for System V
9 ; 8-F F 43be69191228eb3c Unused/padding
10 ; 10-11 G 0200 1 = relocatable, 2 = executable, 3 = shared, 4 = core
11 ; 12-13 H 3e00 Instruction set
12 ; 14-17 I 01000000 ELF Version
13 ; 18-1F J 0400000001000000 Program entry position
14 ; 20-27 K 1c00000000000000 Program header table position - This is actually in th
15 ; 28-2f L 0000000000000000 Section header table position (Don't have one here so
16 ; 30-33 M 01000000 Flags - architecture dependent
17 ; 34-35 N 4000 Header size
18 ; 36-37 O 3800 Size of an entry in the program header table
19 ; 38-39 P 0100 Number of entries in the program header table
20 ; 3A-3B Q 0200 Size of an entry in the section header table
21 ; 3C-3D R b0a9 Number of entries in the section header table [holds n
22 ; 3E-3F S 0f05 Index in section header table with the section name [h
23 ;
24 ; --- Program Header
25 ; OFFSET # Value Purpose
26 ; 1C-1F PA 01000000 Type of segment
27 ;
28 ; 0 = null - ignore the entry

```

28	;		1 = load - clear p_memsz bytes at p_vaddr to 0, then
29	;		2 = dynamic - requires dynamic linking
30	;		3 = interp - contains a file path to an executable t
31	;		4 = note section
32	;	20-23 PB 1c000000	Flags
33	;		1 = executable
34	;		2 = writable
35	;		4 = readable
36	;		In this case the flags are 1c which is 00011100
37	;		The ABI only pays attention to the lowest three bits
38	;	24-2B PC 0000000000000000	The offset in the file that the data for this segment
39	;	2C-33 PD 0000000001000000	Where you should start to put this segment in virtual
40	;	34-3B PE 4000380001000200	Physical Address
41	;	3C-43 PF b0a90f0500000000	Size of the segment in the file (p_filesz) NOTE: Ca
42	;	44-4B PG b0a90f0500000000	Size of the segment in memory (p_memsz) are equa
43	;	4C-43 PH bfaaddee1feebe990	The required alignment for this section (must be a po
44	;		
45	;	Breakdown of the hex dump according to the above data	
46	;	A----- B- C- D- E- F-----	
47	;	00000000 7f 45 4c 46 ba dc fe 21 43 be 69 19 12 28 eb 3c	.ELF...!C.i..(<
48	;		PA-----
49	;	G---- H---- I----- J-----	
50	;	00000010 02 00 3e 00 01 00 00 00 04 00 00 00 01 00 00 00	..>.....
51	;	PB----- PC----- PD-----	
52	;	K----- L-----	
53	;	00000020 1c 00 00 00 00 00 00 00 00 00 00 00 00 00 00
54	;	PD----- PE----- PF-----	
55	;	M----- N---- O---- P---- Q---- R---- S----	
56	;	00000030 01 00 00 00 40 00 38 00 01 00 02 00 b0 a9 0f 05@.8.....
57	;	PF----- PG----- PH-----	
58	;	00000040 00 00 00 00 b0 a9 0f 05 00 00 00 00 bf ad de e1
59	;	PH-----	
60	;	00000050 fe eb e9 90

extendedcomment.asm hosted with ❤ by GitHub

[view raw](#)

The bytes that the instructions are composed of are not contiguous - rather than consisting of a single stream of bytes, some code resides at the beginning and the end of the ELF header with data and `00` bytes in between; in addition to accounting for the output produced by `readElf` above, this may pose a challenge for correct disassembly.

Disassembly with Capstone and Radare2

According to the comments in the source code of the file, the last instruction is located at the last byte of the file - offset `0x53`. Using this information, a simple script to disassemble the code with Capstone can be written:

1	#!/usr/bin/python3
2	
3	from capstone import *
4	
5	CODE = []
6	with open("./bye", "rb") as f:
7	f.seek(4)
8	CODE = f.read()
9	
10	md = md = Cs(CS_ARCH_X86, CS_MODE_64)
11	for i in md.disasm(CODE, 0x1000):
12	print("0x%x:\t%s\t%s" %(i.address, i.mnemonic, i.op_str))

[disassemble_bye_1.py](#) hosted with ❤ by [GitHub](#)
[view raw](#)

This produces the following disassembly:

```
$ ./disassemble_bye.py
0x1000: mov     edx, 0x4321fedc
0x1005: mov     esi, 0x28121969
0x100a: jmp      0x1048
0x100c: add     al, byte ptr [rax]
0x100e: add     byte ptr ds:[rcx], al
0x1011: add     byte ptr [rax], al
0x1013: add     byte ptr [rax + rax], al
0x1016: add     byte ptr [rax], al
0x1018: add     dword ptr [rax], eax
0x101a: add     byte ptr [rax], al
0x101c: sbb     al, 0
0x101e: add     byte ptr [rax], al
0x1020: add     byte ptr [rax], al
0x1022: add     byte ptr [rax], al
0x1024: add     byte ptr [rax], al
0x1026: add     byte ptr [rax], al
0x1028: add     byte ptr [rax], al
0x102a: add     byte ptr [rax], al
0x102c: add     dword ptr [rax], eax
0x102e: add     byte ptr [rax], al
0x1030: add     byte ptr [rax], dil
0x1033: add     byte ptr [rcx], al
0x1035: add     byte ptr [rdx], al
0x1037: add     byte ptr [rax + 0x50fa9], dh
0x103d: add     byte ptr [rax], al
0x103f: add     byte ptr [rax + 0x50fa9], dh
0x1045: add     byte ptr [rax], al
0x1047: add     byte ptr [rdi - 0x11e2153], bh
```

```
0x104d: jmp      0x1038
0x104f: nop
```

This is clearly incorrect. What happened? As it turns out, Capstone is a *linear sweep*-based disassembler (as opposed to *recursive traversal*-based, like radare2)[7][8]. This means that beginning at the start address, it disassembles all bytes as code until the end address, ignoring flow-of-control. In the disassembly above, quite a bit of null bytes and data are being decoded as instructions. We can compensate for this manually somewhat by ignoring the bytes between the `jmp` at offset `0xa` and the `cya` label at offset `0x3c` (see the source code, lines 11 and 27 in particular):

```
1  #!/usr/bin/python3
2
3  from capstone import *
4
5  buf_A = []
6  buf_B = []
7  with open("./bye", "rb") as f:
8      f.seek(4)
9      buf_A = f.read(12)
10     f.seek(0x3c)
11     buf_B = f.read()
12
13  CODE = buf_A + buf_B
14
15  md = md = Cs(CS_ARCH_X86, CS_MODE_64)
16  md.skipdata = True
17  for i in md.disasm(CODE, 0x1000):
18      print("0x%x:\t%s\t%s" %(i.address, i.mnemonic, i.op_str))
```

disassemble_bye_2.py hosted with ♥ by GitHub

[view raw](#)

The disassembly produced after these adjustments is less egregiously erroneous (but still not quite correct):

```
$ ./disassemble_bye_2.py
0x1000: mov     edx, 0x4321fedc
0x1005: mov     esi, 0x28121969
0x100a: jmp      0x1048          <----- jumps beyond the end of the buffer
0x100c: mov     al, 0xa9
0x100e: syscall
0x1010: add     byte ptr [rax], al  <----- error
0x1012: add     byte ptr [rax], al  <----- error
0x1014: mov     al, 0xa9
0x1016: syscall
0x1018: add     byte ptr [rax], al  <----- error
0x101a: add     byte ptr [rax], al  <----- error
0x101c: mov     edi, 0xfed1dead
0x1021: jmp     0x100c
```

```
0x1023: nop
```

At least it somewhat resembles the source code.

How does radare2 fare in disassembling this binary? Not well at all. In fact, it completely fails (maybe I am not using the correct flags?):

```
$ r2 bye
```

```
Warning: Cannot initialize program headers
```

```
Warning: Cannot initialize dynamic strings
```

```
Warning: Cannot initialize dynamic section
```

```
[0x00000004]> pd
```

```
    ;-- entry0:
```

```
    ;-- eip:
```

```
0x00000004      ff      invalid
```

```
0x00000005      ff      invalid
```

```
0x00000006      ff      invalid
```

```
0x00000007      ff      invalid
```

```
0x00000008      ff      invalid
```

```
0x00000009      ff      invalid
```

```
0x0000000a      ff      invalid
```

```
0x0000000b      ff      invalid
```

```
0x0000000c      ff      invalid
```

```
    .
```

```
    .
```

```
    .
```

```
0x00000031      ff      invalid
```

```
0x00000032      ff      invalid
```

```
0x00000033      ff      invalid
```

```
    ;-- section_end.ehdr:
```

```
0x00000034      ff      invalid
```

```
0x00000035      ff      invalid
```

```
0x00000036      ff      invalid
```

```
    .
```

```
    .
```

```
    .
```

```
0x00000040      ff      invalid
```

```
0x00000041      ff      invalid
```

```
0x00000042      ff      invalid
```

```
0x00000043      ff      invalid
```

```
[0x00000004]>
```

Looks like disassembly is not particularly helpful here.

Emulation

Emulation seems to be the most reasonable option. The program responsible for handling emulation of `bye` includes code for handling x86-64 syscalls on lines 26 - 41, allowing us to see the arguments in the registers when the syscall is made:

```
1  #!/usr/bin/python3
2
3  from unicorn import *
4  from unicorn.x86_const import *
5  from capstone import *
6  import struct
7
8
9  BASE = 0x100000
10 STACK_ADDR = 0x0
11 STACK_SIZE = 1024 * 1024
12
13 def read(name):
14     with open(name, 'rb') as f:
15         return f.read()
16
17 #https://github.com/unicorn-engine/unicorn/blob/master/bindings/python/shellcode.py
18 # callback for tracing instructions
19 def hook_code(uc, address, size, user_data):
20     instruction = uc.mem_read(address, size) # read this instruction code from mem
21     md = user_data
22     for i in md.disasm(instruction, address):
23         print(">>> Tracing instruction at 0x%x, instruction size = 0x%x, disassembly:"
24
25
26 def hook_syscall64(mu, user_data):
27     rax = mu.reg_read(UC_X86_REG_RAX)
28     rdi = mu.reg_read(UC_X86_REG_RDI)
29     rsi = mu.reg_read(UC_X86_REG_RSI)
30     rdx = mu.reg_read(UC_X86_REG_RDX)
31
32     print(">>> got SYSCALL with RAX = %d" %(rax))
33
34     if rax == 0xa9: #reboot
35         print(">>> SYSCALL:\treboot\n>>> ARGUMENTS:\trDI = 0x%x\trSI = 0x%x\trDX = 0x"
36             # see syscall table at https://blog.rchapman.org/posts/Linux_System_Call_Table_for_x86_64
37     else:
38         rip = mu.reg_read(UC_X86_REG_RIP)
39         print(">>> Syscall Found at 0x%x: , RAX = 0x%x" %(rip, rax))
40
41 mu.emu_stop()
42
```

```
43 |
44 |
45 | def main():
46 |
47 |     mu = Uc(UC_ARCH_X86, UC_MODE_64)    # initialize emulation engine class
48 |     mu.mem_map(BASE, STACK_SIZE)        # allocate space at base address
49 |     mu.mem_map(STACK_ADDR, STACK_SIZE)   # allocate space for stack
50 |
51 |     mu.mem_write(BASE, read("./linux/bye"))    # write file to memory
52 |     mu.reg_write(UC_X86_REG_RSP, STACK_ADDR + STACK_SIZE - 1) # initialize stack
53 |
54 |     md = Cs(CS_ARCH_X86, CS_MODE_64)    # initialize disassembler engine class
55 |
56 |     # add hooks
57 |     mu.hook_add(UC_HOOK_CODE, hook_code, md)    # pass disassembler engine to hook
58 |     mu.hook_add(UC_HOOK_INSN, hook_syscall64, None, 1, 0, UC_X86_INS_SYSCALL) # hook
59 |
60 |     mu.emu_start(BASE + 0x4, BASE + 0x53)
61 |
62 |     print(">>> Emulation Complete.")
63 |
64 | if __name__ == "__main__":
65 |     main()
```

emulate_bye.py hosted with ❤ by GitHub [view raw](#)

Emulated execution trace:

```
$ ./emulate_bye.py
>>> Tracing instruction at 0x100004, instruction size = 0x5, disassembly:    mov
>>> Tracing instruction at 0x100009, instruction size = 0x5, disassembly:    mov
>>> Tracing instruction at 0x10000e, instruction size = 0x2, disassembly:    jmp
>>> Tracing instruction at 0x10004c, instruction size = 0x5, disassembly:    mov
>>> Tracing instruction at 0x100051, instruction size = 0x2, disassembly:    jmp
>>> Tracing instruction at 0x10003c, instruction size = 0x2, disassembly:    mov
>>> Tracing instruction at 0x10003e, instruction size = 0x2, disassembly:    syscall
>>> got SYSCALL with RAX = 169
>>> SYSCALL:    reboot
>>> ARGUMENTS:  RDI = 0xf00d0000      RSI = 0x28121969      RDX = 0x4321fedc
>>> Emulation Complete.
```

Very nice. Not only do we see the runtime behavior of the program without executing it, but we get essentially correct disassembly as well. According to the source code and the attempt at disassembly using Capstone, the `reboot` syscall is made twice, but obviously only the first one would ever be executed, meaning the instructions following the first `reboot` syscall are unreachable. Perhaps emulation is also useful for analysing obfuscated assembly code? ;)

Conclusion

As we can see, emulation via Unicorn is a very powerful method for analyzing programs that can't be properly parsed or disassembled with the usual tools. However, the difficulty of writing the program that performs the emulation scales with the complexity of the program being emulated. An example of this is the necessity of implementing support manually for interrupts and syscalls. In the next post, somewhat larger programs with a greater range of functionality will be analyzed. Up to this point the start and end addresses of emulation have been manually retrieved from a dump of the target binary; a method of robustly parsing malformed ELF headers will also be explored so that the code start and end offsets can be retrieved in an automated fashion.

Links and References

1. [ELF Crafting: Advanced Anti-analysis techniques for the Linux Platform](#)
2. [Striking Back GDB and IDA debuggers through malformed ELF executables](#)
3. [Screwing elf header for fun and profit](#)
4. [Modern Linux Malware Exposed](#)
5. [Understanding Linux Malware](#)
6. [Linux process execution and the useless ELF header fields](#)
7. [Disassembly of Executable Code Revisited](#) - discusses linear sweep and recursive traversal disassembly algorithms
8. [On Disassembling Obfuscated Assembly](#)

Muppetlabs' Tiny Binaries:

- [The Teensy Files](#)

netspooky's Experiments:

- [source code of "golfclub" binaries on github](#)
- [ELF Binary Mangling Part 1 — Concepts](#)
- [Elf Binary Mangling Pt. 2: Golfin'](#)
- [Elf Binary Mangling Part 3 — Weaponization](#)

Unicorn Engine materials:

- [Unicorn Engine tutorial](#)
- [Unicorn Engine Reference \(Unofficial\)](#)
- [sample_x86.py](#)
- [shellcode.py](#)

SHARE    

TAGS



PREVIOUS POST

[Inserting Debugging Instrumentation into an Internal Function Using Redirect-to-PLT](#)

♥ Recommend

🐦 Tweet

f Share

Sort by Oldest ▾



Join the discussion...

LOG IN WITH

OR SIGN UP WITH DISQUS (?)



Name



Jim • 3 days ago

I'm not seeing the scripts/excerpts you mention, am I missing something or are those scripts available so I can recreate your results?

^ | ▾ • Reply • Share ›

✉ Subscribe

D Add Disqus to your site

🔒 Disqus' Privacy Policy

DISQUS

content created by julian

